

Agenda

1. Spark Streaming
2. Assignment 9 Overview
3. Demo

Streaming with Spark

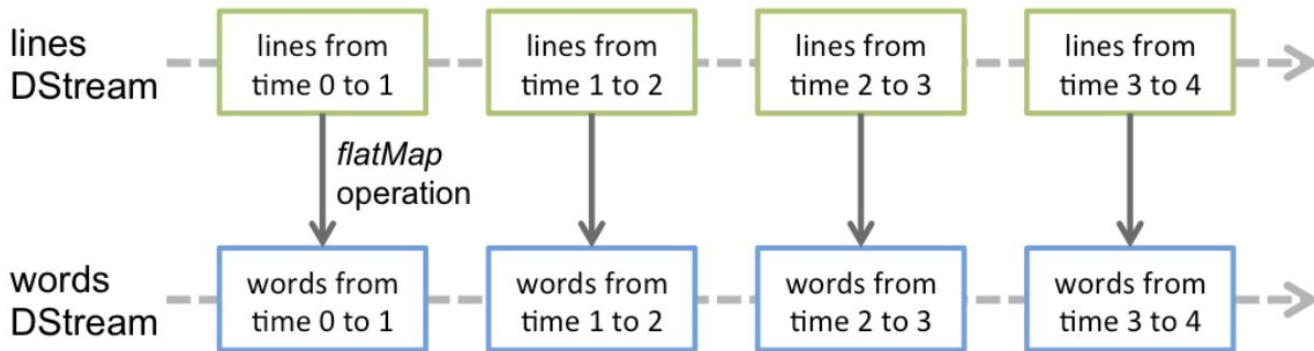


<https://spark.apache.org/docs/latest/streaming-programming-guide.html>


Dstream - discretized stream



Wordcount example



Windows operations



- *window length* - The duration of the window.
- *sliding interval* - The interval at which the window operation is performed.

Python env setup

```
from pyspark import SparkContext, SparkConf  
from pyspark.streaming import StreamingContext  
from pyspark.streaming.kafka import KafkaUtils
```

```
--packages org.apache.spark:spark-streaming-kafka-0-8_2.11:2.0.0
```

Advanced Sources

Python API As of Spark 2.4.4, out of these sources, Kafka, Kinesis and Flume are available in the Python API.

Java env setup

1. Sample maven config for Kafka Producer

```
<dependency>  
    <groupId>org.apache.kafka</groupId>  
    <artifactId>kafka-clients</artifactId>  
    <version>2.0.0</version>  
</dependency>
```

2. Sample maven config for Spark Job

```
<dependency>  
    <groupId>org.apache.spark</groupId>  
    <artifactId>spark-streaming_2.12</artifactId>  
    <version>2.4.4</version>  
    <scope>provided</scope>  
</dependency>  
<dependency>  
    <groupId>org.apache.spark</groupId>  
    <artifactId>spark-streaming-kafka-0-10_2.12</artifactId>  
    <version>2.4.4</version>  
</dependency>
```

3. Either generate the JAR with all dependencies or use --package option to define the dependency

Kafka Connector

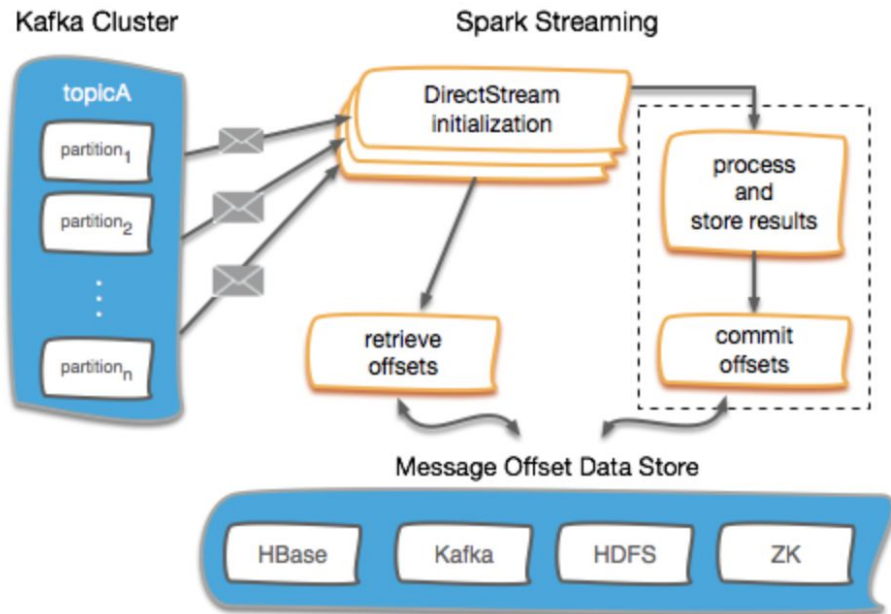
The Kafka project introduced a new consumer API between versions 0.8 and 0.10, so there are 2 separate corresponding Spark Streaming packages available. Please choose the correct package for your brokers and desired features; note that the 0.8 integration is compatible with later 0.9 and 0.10 brokers, but the 0.10 integration is not compatible with earlier brokers.

Note: Kafka 0.8 support is deprecated as of Spark 2.3.0.

| | spark-streaming-kafka-0-8 | spark-streaming-kafka-0-10 |
|----------------------------|---------------------------|----------------------------|
| Broker Version | 0.8.2.1 or higher | 0.10.0 or higher |
| API Maturity | Deprecated | Stable |
| Language Support | Scala, Java, Python | Scala, Java |
| Receiver DStream | Yes | No |
| Direct DStream | Yes | Yes |
| SSL / TLS Support | No | Yes |
| Offset Commit API | No | Yes |
| Dynamic Topic Subscription | No | Yes |

<https://spark.apache.org/docs/latest/streaming-kafka-integration.htm>

Offsets management - important for Prod deployment



[figure 1 – high-level flow for managing offsets]

Kafka and Spark Streaming Demo

SetupKafka-AWS-PublicSubnets.json - Cloudformation script to setup Kafka in AWS.

Useful commands

Run java spark job

```
spark-submit --class edu.harvard.e88.lab9spark.SparkStreamConsumer lab9spark-0.0.1-SNAPSHOT-jar-with-dependencies.jar  
mytopic kafkabrokerIP:9092
```

Run Python spark job

```
spark-submit --packages org.apache.spark:spark-streaming-kafka-0-8_2.11:2.0.0 SparkStreamingConsumer.py kafkabrokerIP:9092  
mytopic
```

Run Kafka Producer

```
java -cp lab9-0.0.1-SNAPSHOT-jar-with-dependencies.jar edu.harvard.e88.lab9.WeblogsKafkaProducer mytopic kafkabrokerIP:9092
```

Create Kafka Topic

```
./kafka-topics.sh --zookeeper ZookeeperIP:2181 --create --topic mytopic --partitions 2 --replication-factor 1
```

UpdateStateByKey

Additional References

<https://blog.cloudera.com/offset-management-for-apache-kafka-with-apache-spark-streaming/>

<https://aws.amazon.com/blogs/big-data/real-time-stream-processing-using-apache-spark-streaming-and-apache-kafka-on-aws/>