


Lab 11: Cassandra

Agenda:

1. HW review
2. Local Cassandra setup
3. working with Cassandra from Java
4. working with Cassandra from Python
5. Cassandra tools
6. Cassandra cluster setup with Docker

Cassandra:

CAP Theorem



1. Fast Distributed (Column Family NoSQL)
2. Database High availability Linear Scalability High Performance
3. Fault tolerant
4. Multi-Data Center Support
5. Easy to operate

Cassandra Data Modeling

1. Access Pattern data modeling
2. Spread data evenly across the cluster
3. Use denormalization

Local Cassandra setup

Download from: <http://cassandra.apache.org/download/>

Start: `./bin/cassandra`

CQLSH: `./bin/cqlsh localhost`

Stop Cassandra process on tarball installation:

Find the Cassandra Java process ID (PID), and then kill the process using its PID number:

```
$ ps auwx | grep CassandraDaemon
```

```
$ kill pid
```

Cassandra with Python:

pip install cassandra-driver

Use driver:

http://datastax.github.io/python-driver/getting_started.html#asynchronous-queries

Cassandra tools

`./bin/nodetool`

`./bin/nodetool info` - quick summary of the node

./bin/nodetool version - version of cassandra

./bin/nodetool status <keyspace_name> - status of cluster/keyspace

./bin/nodetool cfstats <keyspace_name.table_name> - info on the table

```
CREATE KEYSPACE lab10 WITH REPLICATION = { 'class' : 'SimpleStrategy',  
'replication_factor' : 1 };
```

```
CREATE TABLE IF NOT EXISTS lab10.sensor_metrics (  
    sensor_id UUID,  
    time_hour timestamp,  
    sensor_type text,  
    reading_time timestamp,  
    metric float,  
    PRIMARY KEY ((sensor_type, time_hour), reading_time)  
) WITH CLUSTERING ORDER BY (reading_time asc);
```

```
INSERT INTO lab10.sensor_metrics  
(sensor_type, time_hour, reading_time, sensor_id, metric)  
VALUES ('p_type1', '2019-11-11 21:00:00', '2019-11-11 21:05:00', uuid(), 1.0);
```

```
SELECT * from lab10.sensor_metrics WHERE sensor_type='p_type1' AND time_hour =  
'2019-11-11 21:00:00';
```

Docker setup for Cassandra cluster

```
docker run --name cassandra-node1 -d -p 9042:9042 -e  
CASSANDRA_CLUSTER_NAME=Lab10 -e  
CASSANDRA_ENDPOINT_SNITCH=GossipingPropertyFileSnitch -e  
CASSANDRA_DC=datacenter1 cassandra
```

```
docker inspect --format '{{ .NetworkSettings.IPAddress }}' cassandra-node1
```

```
docker run --name cassandra-node2 -d -e CASSANDRA_SEEDS="$(docker inspect
--format='{{ .NetworkSettings.IPAddress }}' cassandra-node1)" -e
CASSANDRA_CLUSTER_NAME=Lab10 -e
CASSANDRA_ENDPOINT_SNITCH=GossipingPropertyFileSnitch -e
CASSANDRA_DC=datacenter1 cassandra
```

Execute below statement to check nodetool status to check the status of cluster . Don't proceed to node3 untill you see second node in UN

```
docker exec cassandra-node1 nodetool status
```

```
docker run --name cassandra-node3 -d -e CASSANDRA_SEEDS="$(docker inspect
--format='{{ .NetworkSettings.IPAddress }}' cassandra-node1)" -e
CASSANDRA_CLUSTER_NAME=Lab10 -e
CASSANDRA_ENDPOINT_SNITCH=GossipingPropertyFileSnitch -e
CASSANDRA_DC=datacenter2 cassandra
```

Execute below statement to check nodetool status to check the status of cluster .Don't proceed until you see all three nodes in UN

```
docker exec cassandra-node1 nodetool status
```

Execute below statement to connect to node-1 and connect to cqlsh session

```
docker exec -it cassandra-node1 cqlsh
```

```
CREATE KEYSPACE lab10
WITH replication = {
    'class' : 'NetworkTopologyStrategy',
    'datacenter1' : 2,
    'datacenter2' : 1
};
```

```
CREATE TABLE lab10.person (
    id int primary key,
    name text
);
```

```
docker stop cassandra-node1 cassandra-node2 cassandra-node3
```

```
CREATE KEYSPACE lab10 WITH REPLICATION = { 'class' : 'SimpleStrategy',  
'replication_factor' : 1 };
```

```
CREATE TABLE IF NOT EXISTS lab10.sensor_metrics (  
    sensor_id UUID,  
    time_hour timestamp,  
    sensor_type text,  
    reading_time timestamp,  
    metric float,  
    PRIMARY KEY ((sensor_type, time_hour), reading_time)  
) WITH CLUSTERING ORDER BY (reading_time asc);
```

```
INSERT INTO lab10.sensor_metrics  
    (sensor_type, time_hour, reading_time, sensor_id, metric)  
VALUES ('p_type1', '2019-11-11 21:00:00', '2019-11-11 21:05:00', uuid(), 1.0);
```

```
SELECT * from lab10.sensor_metrics WHERE sensor_type='p_type1' AND time_hour =  
'2019-11-11 21:00:00';
```