

# Lab12: ElasticSearch Basics

by Andriy Pyshchuk

# Setup Elasticsearch + Kibana

1. Download Elasticsearch <https://www.elastic.co/downloads/elasticsearch>
2. Download Kibana <https://www.elastic.co/downloads/kibana>
3. Unpack

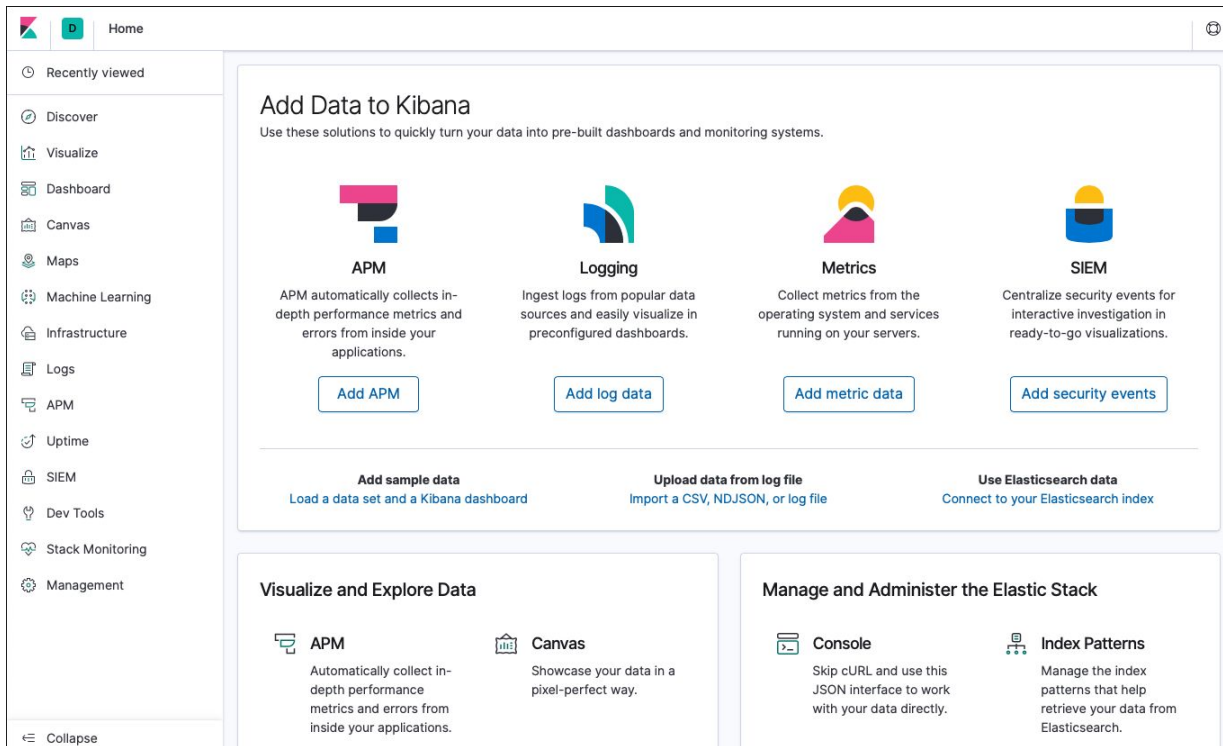
# Run Elasticsearch

1. Run Elasticsearch bin/elasticsearch
2. Open localhost:9200

```
{
  "name" : "zTJbz81",
  "cluster_name" : "elasticsearch",
  "cluster_uuid" : "ZMpUUNdIRZW6xtyycodIDg",
  "version" : {
    "number" : "6.2.3",
    "build_hash" : "c59ff00",
    "build_date" : "2018-03-13T10:06:29.741383Z",
    "build_snapshot" : false,
    "lucene_version" : "7.2.1",
    "minimum_wire_compatibility_version" : "5.6.0",
    "minimum_index_compatibility_version" : "5.0.0"
  },
  "tagline" : "You Know, for Search"
}
```

# Run Kibana

1. Run Kibana bin/kibana
2. Open localhost:5601



# Terminology

## Relation Databases

- Database
- Table
- Row
- Column
- Schema



## Elasticsearch

Index

Type

Document

Fields

Mapping



# Document sample

```
{
  "_index": "newuser",
  "_type": "user",
  "_id": "AV_RUs33bC0qJzkoLfN1",
  "_score": 1,
  "_source": {
    "name": "Andriy",
    "age": "30"
  }
}
```

```

  "_index": "nyc_restaurants",
  "_type": "inspection",
  "_id": "113488",
  "_score": 1,
  "_source": {
    "Dbid": "LA MIA PIZZA",
    "Inspection_Type": "Cycle Inspection / Initial Inspection",
    "Inspection_Date": [
      "2014-07-02T00:00:00"
    ],
    "Action": "Violations were cited in the following area(s).",
    "Violation_Code": "02G",
    "Score": 26,
    "Building": "1488      ",
    "Grade_Date": null,
    "Critical_Flag": "Critical",
    "Camis": 41030858,
    "Zipcode": 10075,
    "Violation_Description": "Cold food item held above 41° F (smoked fish and reduced oxygen packaged foods above 38 °F) except during necessary preparation.",
    "Phone": "2124721200",
    "Cuisine_Description": "Pizza",
    "Grade": "",
    "Street": "1 AVENUE",
    "Coord": [
      -73.9531563,
      40.77127040000001
    ],
    "Record_Date": "2016-03-21T00:00:00",
    "Address": "1488 1 AVENUE MANHATTAN,NY",
    "Boro": "MANHATTAN"
  },
}
```

# Mapping sample

PUT my\_index

```
{
  "mappings": {
    "my_type": {
      "properties": {
        "full_text": {
          "type": "text"
        },
        "exact_value": {
          "type": "keyword"
        }
      }
    }
  }
}
```

# ES field types

1. Basic data types (String, numeric, date, boolean, etc)
2. Geo data types (Geo-point, Geo-shape)
3. Specialized data types (IP, completion, join, alias, etc)



# Analyzed vs Not Analyzed

**Input:** Hello Class 123

## **Analyzed (text)**

**Token1:** hello **Token2:** class **Token3:** 123

**Allows to search like:** hello class, hello 123, etc

**Pros** – easy to find something in text

**Cons** – consumes more resources

**Input:** Hello Class 123

## **Not Analyzed (keyword)**

**Token1:** Hello Class 123

**Allows to search only:** Hello Class 123

**Pros** – consumes less resources, very effective for search in log files

**Cons** – you should exactly know what you are searching for

# Analyzed vs Not Analyzed use cases

## Analyzed (text)

**example:** titles that contain the word "jobs".  
query:"title:jobs".

**doc1** : title:developer jobs in boston

**doc2** : title:java coder jobs in vancouver

**doc3** : title:unix designer jobs in austin

**doc4** : title:database manager vacancies in  
montreal

this is going to retrieve title1, title2, title3

## Not Analyzed (keyword)

**example:** get all the logs from machine 1.  
query:"workstation:machine 1".

**doc1:** workstation:machine 1, log: somestring

**doc2:** workstation:machine 2, log: somestring

**doc3:** workstation:machine 1, log: somestring

**doc4:** workstation:machine 4, log: somestring

This is going to retrieve results from doc1 and  
doc3 only

# Demo

creating index, mapping, shards explanation

# Add/update document

```
POST index1/my_type
{
  "name" : "Taras",
  "age" : 25
}
```

```
PUT index1/my_type/2
{
  "name" : "Taras",
  "age" : 25
}
```

```
{
  "_index": "index1",
  "_type": "my_type",
  "_id": "2",
  "_version": 1,
  "result": "created",
  "_shards": {
    "total": 2,
    "successful": 1,
    "failed": 0
  },
  "_seq_no": 1,
  "_primary_term": 1
}
```

# Range search

Add couple documents:

POST index1/my\_type

```
{  
  "name" : "Bob",  
  "age" : "70"
```

```
}
```

POST index1/my\_type

```
{  
  "name" : "Andriy",  
  "age" : "30"  
}
```

## Query

GET /textsearch/\_search

```
{ "query": {  
  "range" : {  
    "age" : { "gte" : 30,  
              "lte" : 69 }  
  }  
}
```

# Wildcard search

Add couple documents:

**POST** index1/my\_type

```
{  
  "name" : "Bob",  
  "age" : "70"  
}
```

**POST** index1/my\_type

```
{  
  "name" : "Andriy",  
  "age" : "30"  
}
```

**Query**

**GET** /textsearch/\_search

```
{ "query": {  
  "wildcard" : {  
    "name" : "b*"  
  }  
}
```

# Fuzzy search

Add couple documents:

**POST** students/student

```
{  
  "name" : "Bob",  
  "age" : "20"  
}
```

**POST** students/student

```
{  
  "name" : "Andriy",  
  "age" : "30"  
}
```

```
GET students/_search  
{  
  "query": {  
    "fuzzy" : {  
      "name" : {  
        "value": "Bim",  
        "fuzziness": 2  
      }  
    }  
  }  
}
```

# Exists Query

Returns documents that have at least one non-null value in the original field:

```
GET textsearch/_search
{
  "query": {
    "exists" : { "field" : "name" }
  }
}
```



# Doesn't exists Query

Returns documents that have at least one non-null value in the original field:

```
GET textsearch/_search
{
  "query": {
    "bool": {
      "must_not": {
        "exists" : { "field" : "name" }
      }
    }
  }
}
```

# Aggregation (Top)

GET students/\_search?size=0

```
{
  "size": 0,
  "aggregations": {
    "2": {
      "terms": {
        "field": "name",
        "size": 5,
        "order": {
          "_count": "desc"
        }
      }
    }
  }
}
```

---

```
"aggregations": {
  "2": {
    "doc_count_error_upper_bound": 0,
    "sum_other_doc_count": 0,
    "buckets": [
      {
        "key": "Alex",
        "doc_count": 2
      },
      {
        "key": "Andriy",
        "doc_count": 1
      },
      {
        "key": "Bob",
        "doc_count": 1
      },
      {
        "key": "Yuriy",
        "doc_count": 1
      }
    ]
  }
}
```

# Aggregation (cardinality)

POST students/\_search?size=0

```
{
  "aggregations" : {
    "type_count" : {
      "cardinality" : {
        "field" : "name"
      }
    }
  }
}
```

```
{
  "took": 0,
  "timed_out": false,
  "_shards": {
    "total": 5,
    "successful": 5,
    "skipped": 0,
    "failed": 0
  },
  "hits": {
    "total": 5,
    "max_score": 0,
    "hits": []
  },
  "aggregations": {
    "type_count": {
      "value": 4
    }
  }
}
```

# Bulk add

POST /museums\_index/doc/\_bulk?refresh

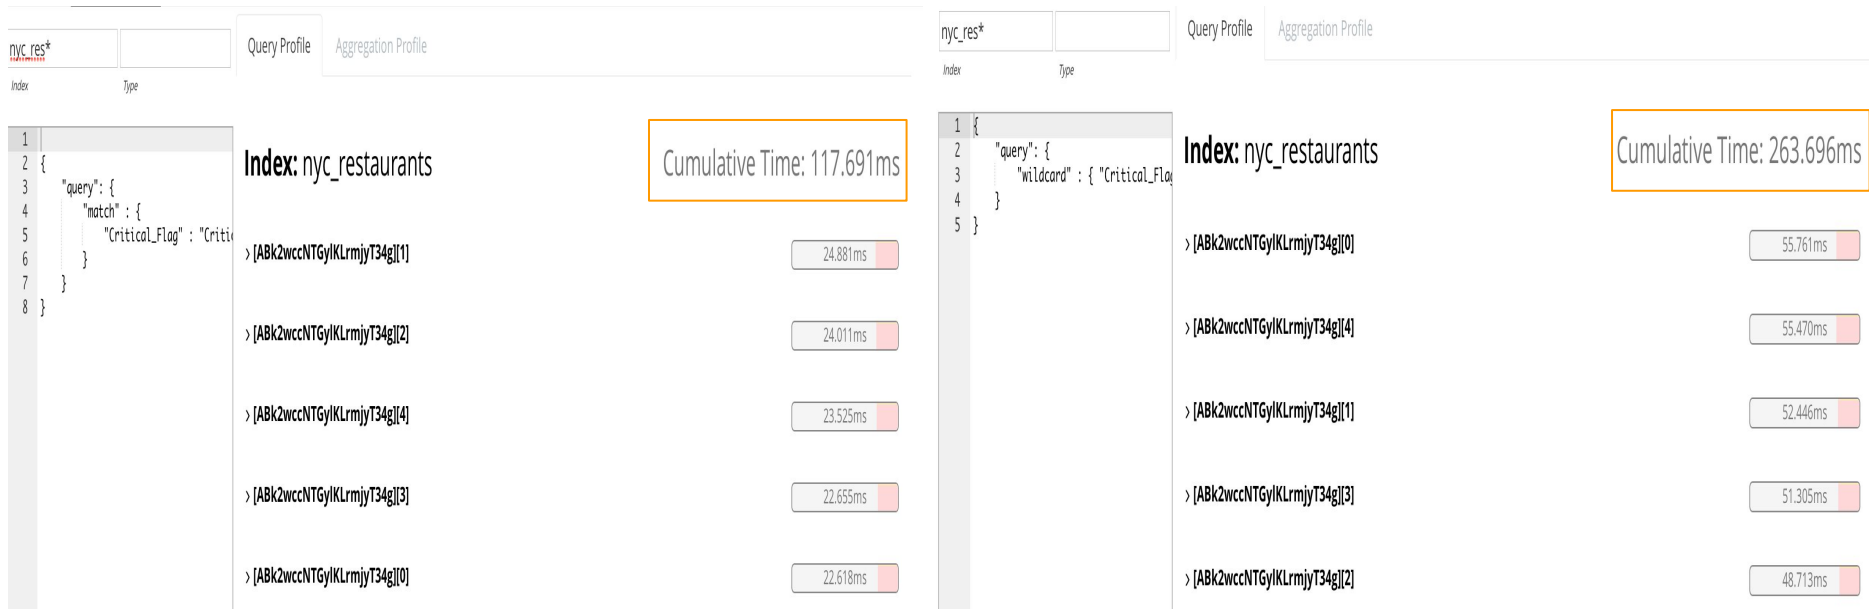
```
{"index":{"_id":1}}
{"location": "52.374081,4.912350", "name":
"NEMO Science Museum"}
{"index":{"_id":2}}
{"location": "52.369219,4.901618", "name":
"Museum Het Rembrandthuis"}
{"index":{"_id":3}}
{"location": "52.371667,4.914722", "name":
"Nederlands Scheepvaartmuseum"}
{"index":{"_id":4}}
{"location": "51.222900,4.405200", "name":
"Letterenhuis"}
{"index":{"_id":5}}
{"location": "48.861111,2.336389", "name":
"Musée du Louvre"}
{"index":{"_id":6}}
{"location": "48.860000,2.327000", "name":
"Musée d'Orsay"}
```

```
{
  "took": 263,
  "errors": false,
  "items": [
    {
      "index": {
        "_index": "museums_index",
        "_type": "doc",
        "_id": "1",
        "_version": 1,
        "result": "created",
        "forced_refresh": true,
        "_shards": {
          "total": 2,
          "successful": 1,
          "failed": 0
        },
        "created": true,
        "status": 201
      }
    },
    {
      "index": {
        "_index": "museums_index",
        "_type": "doc",
        "_id": "2",
        "_version": 1,
```

# NYC Restaurant Inspection Data structure

t	Action	Q Q □ *	Violations were cited in the following area(s).
t	Address	Q Q □ *	18704 HILLSIDE AVENUE QUEENS,NY
t	Boro	Q Q □ *	QUEENS
t	Building	Q Q □ *	18704
#	Camis	Q Q □ *	40,853,123
📍	Coord	Q Q □ *	-73.7125054, 40.7364223
t	Critical_Flag	Q Q □ *	Critical
t	Cuisine_Description	Q Q □ *	American
t	Dbn	Q Q □ *	LA VERITE RESTAURANT BAR
t	Grade	Q Q □ *	A
🕒	Grade_Date	Q Q □ *	March 19th 2016, 02:00:00.000
🕒	Inspection_Date	Q Q □ *	March 19th 2016, 02:00:00.000
t	Inspection_Type	Q Q □ *	Cycle Inspection / Initial Inspection
t	Phone	Q Q □ *	7184547479
🕒	Record_Date	Q Q □ *	March 21st 2016, 02:00:00.000
#	Score	Q Q □ *	7
t	Street	Q Q □ *	HILLSIDE AVENUE
t	Violation_Code	Q Q □ *	06F
t	Violation_Description	Q Q □ *	Wiping cloths soiled or not stored in sanitizing solution.
#	Zipcode	Q Q □ *	11,432
t	_id	Q Q □ *	81544
t	_index	Q Q □ *	nyc_restaurants
#	_score	Q Q □ *	-
t	_type	Q Q □ *	inspection

# Wildcard search - profiling



# Cardinality aggregation

A single-value metrics aggregation that calculates an approximate count of distinct values.

```
GET /nyc_res*/_search?size=0
```

```
{
  "aggs" : {
    "type_count" : {
      "cardinality" : {
        "field" : "Zipcode"
      }
    }
  }
}
```

```
GET /nyc_res*/_search?size=0
```

```
{
  "aggs" : {
    "type_count" : {
      "cardinality" : {
        "field" : "Cuisine_Description"
      }
    }
  }
}
```

# Java connect to Elasticsearch (versions 6+)

```
public void init() throws Exception {
    logger.info("Initializing ElasticSearchClient ...");
    // connect to elasticsearch cluster
    Settings settings = Settings.builder().put(CLUSTER_NAME, esClusterName).build();
    try {
        esTransportClient = new PreBuiltTransportClient(settings);
        esTransportClient.addTransportAddress(
            new TransportAddress(InetAddress.getByName("localhost"), 9300));
    } catch (Exception e) {
        logger.error("Exception trying to connect and create ElasticSearch Client: " + e.getMessage());
        throw e;
    }
}
```



# Python application

```
from datetime import datetime
from elasticsearch import Elasticsearch
es = Elasticsearch()

doc = {
    'author': 'kimchy',
    'text': 'Elasticsearch: cool. bonsai cool.',
    'timestamp': datetime.now(),
}

res = es.index(index="test-index", doc_type='tweet', id=1, body=doc)
print(res['result'])

res = es.get(index="test-index", doc_type='tweet', id=1)
print(res['_source'])

es.indices.refresh(index="test-index")

res = es.search(index="test-index", body={"query": {"match_all": {}}})
print("Got %d Hits:" % res['hits']['total'])

for hit in res['hits']['hits']:
    print("%(timestamp)s %(author)s: %(text)s" % hit["_source"])
```

apysh@macbook325:~\$ sudo python es.py

created

{u'text': u'Elasticsearch: cool. bonsai cool.', u'author': u'kimchy', u'timestamp':  
u'2018-11-18T17:16:04.563802'}

Got 1 Hits:

2018-11-18T17:16:04.563802 kimchy: Elasticsearch: cool. bonsai cool.

# Demo

Visualizations and Dashboards

# Useful commands and links

**{host}:9200/\_cluster/health** - shows cluster health

**{host}:9200/\_cluster/stats** - shows cluster stats

**{host}:9200/\_cat** - cat api, shows all possible options with cat api

**{host}:9200/\_cat/indices** - shows all indices

**Add ?v in the end of call** - to have headers for table output (\_cat api)

**Add ?pretty** - to have nice json output (cluster health)

<https://www.elastic.co/guide/en/elasticsearch/reference/current/index.html> - ES documentation with good samples

[https://github.com/elastic/examples/tree/master/Exploring%20Public%20Datasets/nyc\\_restaurants](https://github.com/elastic/examples/tree/master/Exploring%20Public%20Datasets/nyc_restaurants) - New York City restaurant inspection data